

WA Precipitation from 1895-2017 Report

Initially we import WA precipitation data with date from 1895 to 2017 and monthly average precipitation (inches).

```
library('ggplot2')
library('forecast')

## Warning: package 'forecast' was built under R version 3.4.4

library('tseries')

## Warning: package 'tseries' was built under R version 3.4.4

library('lubridate')

##
## Attaching package: 'lubridate'
## The following object is masked from 'package:base':
##
##      date


data = read.csv('https://www.ncdc.noaa.gov/cag/statewide/time-series/45-pcp-all-1-1895-2017.csv', header=1)
date <- data$Washington[-c(1,2,3)]
value <- data$Precipitation[-c(1,2,3)]
df <- data.frame(date, value)
df$date <- ymd(paste0(as.character(df$date), '01'))
head(df)

##           date value
## 1 1895-01-01  7.53
## 2 1895-02-01  2.54
## 3 1895-03-01  2.95
## 4 1895-04-01  3.00
## 5 1895-05-01  3.29
## 6 1895-06-01  0.62
```

We then use `ts()` function to convert the the numeric vector, “precipitation”, to time series object. After that, we plot the time series of out raw data. We then use `tsclean()` to identify and replace outliers using series smoothing and decomposition. However, the cleaned version seems not any different. (Probably because no outliers).


```
value_ts = ts(df[, c('value')])
df$clean_value = tsclean(value_ts)
ggplot() +
  geom_line(data = df, aes(x = date, y = value_ts, color = "raw data" )) +
  geom_line(data = df, aes(x = date, y = clean_value, color = "cleaned data")) + ylab('WA precipitation')

## Don't know how to automatically pick scale for object of type ts. Defaulting to continuous.
```



Since we use monthly data and precipitation is related to season, we might consider to decompose the seasonal component.

```
value_ma <- ts(na.omit(df$clean_value), frequency=12)
decomp <- stl(value_ma, s.window="periodic")
deseasonal_value <- seasadj(decomp)
autoplot(decomp)
```



Then we use Dickey-Fuller test to determine whether the time series is stationary.

```
adf.test(value_ma, alternative = "stationary")
```

```
## Warning in adf.test(value_ma, alternative = "stationary"): p-value smaller
## than printed p-value
```

```
##
## Augmented Dickey-Fuller Test
##
## data: value_ma
## Dickey-Fuller = -9.2022, Lag order = 11, p-value = 0.01
## alternative hypothesis: stationary
```

Our null hypothesis is “the time series is not stationary”. Notice the p-value is less than 0.05, we can reject our null hypothesis. Thus the time series is stationary.

We will run the `auto.arima` function to see which ARIMA model fits our data best.


```
fit <- auto.arima(deseasonal_value, seasonal=FALSE)
fit
```

```
## Series: deseasonal_value
## ARIMA(2,0,0) with non-zero mean
##
## Coefficients:
##      ar1      ar2      mean
##      0.079  0.041  300.4608
## s.e.  0.026  0.026   3.5287
##
## sigma^2 estimated as 14268: log likelihood=-9152.4
```

```
## AIC=18312.8   AICc=18312.83   BIC=18333.99
```

Finally we can use this model to do the forecast.

```
fcast <- forecast(fit)
autoplot(fcast)
```



```
fitwithseason <- auto.arima(deseasonal_value, seasonal=TRUE)
fcast2 <- forecast(fitwithseason)
autoplot(fcast2)
```

